

SMART-FACILITY INCIDENT DETECTION · ASSESSMENT PROTOTYPE

Factory Safety Sentinel

A predictive, multimodal industrial-safety system — physics- and ML-driven gas-risk forecasting fused with real-time computer-vision PPE and zone supervision, wired into one explainable incident and review workflow.

PHYSICS + XGBOOST + GRU

YOLO11N + BYTETRACK

FULL-STACK, TESTED, EVIDENCE-BACKED

Prepared for the Smart-Facility Incident Detection assessment · Contact: mohammad.2003.1m@gmail.com
Prototype — not certified for real industrial safety decisions.

NO CROSSING

Gas exposure and crossing predicted within
10m radius of worker + CPU

PEOPLE AT RISK

0

confirmed in gas zone

Camera — camera-1

OK

Worker helmet: UNKNOWN · vest: UNKNOWN · gas zone:
#46 · OUTSIDE · overhead zone: OUTSIDE

Worker helmet: UNKNOWN · vest: UNKNOWN · gas zone:
#7 · UNKNOWN · overhead zone: UNKNOWN

Worker helmet: UNKNOWN · vest: UNKNOWN · gas zone:
#2 · OUTSIDE · overhead zone: OUTSIDE

Worker helmet: UNKNOWN · vest: UNKNOWN · gas zone:
#2 · OUTSIDE · overhead zone: OUTSIDE

Worker helmet: UNKNOWN · vest: UNKNOWN · gas zone:
#2 · OUTSIDE · overhead zone: OUTSIDE

source: CV_MODEL · model: ppe-yolo11n-1.1 · last frame: 0s · fps: 8.7
configured zones (v1.0): Gas exposure zone · Overhead work zone ·
Mandatory-vest zone · Restricted zone


Active

Resolved

System is healthy.

Industrial safety monitoring is fragmented and reactive

- **Gas leaks develop gradually.** A slow CO2 buildup can cross a dangerous threshold before any human is watching the reading.
- **PPE violations expose workers to preventable risk** — a missing helmet or vest is only caught if someone happens to look.
- **Restricted areas need continuous supervision** that a manager cannot realistically provide across a whole floor, every shift.
- **Sensors, cameras, and incident logs are usually separate systems** — nothing correlates a gas trend with a camera event with a reviewable record.



The real, working dashboard: one screen, live gas trend + forecast + camera + incidents — replacing the fragmented view above.

MVP scope vs. certified future deployment


DELIVERED IN THIS MVP

- Sensor-based risk monitoring (seeded, replayable CO2 simulator)
- Predictive analysis (physics mass-balance + GRU residual + XGBoost leak classifier)
- Computer vision (YOLO11n + ByteTrack PPE/zone supervision)
- Working manager dashboard with live gas/camera/incident views
- Fully functioning backend (FastAPI, SQLite, WebSocket, REST)
- Simulation / digital twin for testing without a physical factory
- Evidence capture + full incident review/audit workflow

DEFERRED TO CERTIFIED DEPLOYMENT

- Real factory sensor/camera hardware integration
- Multi-camera, multi-zone, multi-worker at production scale
- Regulatory certification for real safety-critical decisions
- Domain-adapted vision model trained on real factory footage
- Production security hardening (auth, encrypted evidence storage)

End-to-end architecture




One ingestion contract for both paths (CLAUDE.md invariant #1): simulator readings and real camera/replay frames flow through the same versioned pipeline into the same risk/incident engine — a real device could be swapped in without changing the backend.

Three trained ML models, one analytical model, three deterministic components

Component	Type	Role
XGBoost leak classifier	Trained ML	Calibrated leak-probability estimate from sliding-window sensor features
Residual GRU forecast model	Trained ML	Learns the physics model's residual error to sharpen the 60-min CO2 forecast
YOLO11n PPE detector	Trained CV Model	Detects person / helmet / vest / no_helmet on camera frames
Physics mass-balance model	Analytical, Not ML	First-principles gas concentration ODE — the safe fallback when ML is unavailable
ByteTrack	Tracking Algorithm	Associates detections into anonymous, session-local worker tracks — not trained
Restricted-zone polygons	Deterministic Geometry	Configured point-in-polygon regions, not a learned "zone" class
Risk / severity policy	Deterministic Logic	Versioned rule table converting evidence into severity + reason codes

Proactive prediction, not threshold-only reaction



Every stage runs on **every tick**, using the last 10 simulated hours of history. The forecast projects 60 minutes ahead — so a manager sees a predicted threshold crossing *before* it happens, with an explicit "Time-to-Action" estimate, not just a current-value alarm.

Time-to-Action: the equations behind the forecast

WELL-MIXED MASS BALANCE

$$dC/dt = Q/V \cdot (C_{in} - C) + G/V$$

STEADY-STATE CONCENTRATION

$$C_{ss} = C_{in} + G/Q$$

TIME CONSTANT

$$\tau = V/Q$$

CONCENTRATION RESPONSE

$$C(t) = C_{ss} + (C_0 - C_{ss}) \cdot e^{(-t/\tau)}$$

TIME-TO-ACTION

$$t = -\tau \cdot \ln((C_{thresh} - C_{ss}) / (C_0 - C_{ss}))$$

WHAT THIS MEANS FOR A MANAGER

Given the zone's volume, ventilation rate, and current emission source, the system solves exactly when CO2 will cross the NIOSH action reference — and reports it in plain minutes, e.g. *"CO2 may reach the 5000 ppm action reference in 34 minutes."*

Deliberately called **Time-to-Action**, never "time to harm" — the 5000 ppm reference is an occupational action threshold, not an immediate-harm line (CLAUDE.md).

Forecast accuracy and leak-classifier quality

MATCHED PHYSICS-VS-HYBRID BENCHMARK (N=1092)

57.19_{ppm}

47.59_{ppm}

PHYSICS-ONLY MAE

HYBRID PHYSICS+GRU MAE

▼ 16.8% ERROR REDUCTION

Identical held-out scenario/point set for both — a fair, matched comparison (models/evaluation/gru_benchmark_report.json).

LEAK CLASSIFIER (XGBOOST, CALIBRATED)

0.965

0.022

0.942

PR-AUC

BRIER (CALIBRATED)


F1

n=900 held-out samples, 198 positive. Forecast inference latency: 1.90ms median / 2.41ms p95.

Not the same number: a separate, broader physics-only evaluation across 10 scenarios / 120 points reports 95.50 ppm MAE — a different, larger evaluation set. Never compared directly against the 47.59 ppm matched hybrid figure above.

YOLO11n + ByteTrack, real evidence, no fabricated CV

- **YOLO11n** — active version **1.1**, 640px input, fine-tuned from COCO-pretrained weights
- **Class schema:** person · helmet · vest · no_helmet — the only 4 classes exposed at runtime
- **ByteTrack** assigns anonymous, session-local track IDs — no facial recognition, ever
- **PPE dwell:** 3s persistence required before a violation, 5s clear before resolving
- **Restricted-zone foot-point logic:** bottom-center of the person box tested against a configured polygon, with a 2s enter / 2s exit debounce
- **Real evidence capture:** the actual annotated camera frame that triggered the incident is saved — never a placeholder



Genuine annotated frame from the final interview-demo video: real person/PPE boxes, confidence scores, zone overlay, and burned-in model version + timestamp.

v1.1 active vs. v1.2 candidate — an honest promotion decision

WHY MORE DATA WAS TRIED

v1.1's no_helmet recall (0.175) was the known weak point. A second, independent dataset — **Industrial Safety** (Roboflow Universe, MIT license, 28,119 images) — was acquired to fine-tune a candidate.

Canonical mapping: hardhat → helmet, no_hardhat → no_helmet, person → person, safety_vest → vest. Deduplicated (3 exact dupes dropped) and leakage-checked before a class-balanced 8,000-image training subset was selected — augmented/near-duplicate video frames were capped, never counted as unique scenes.

PROMOTION DECISION: REJECTED

The candidate completed only 7 of a planned 50 epochs (externally interrupted for runtime, not early-stopped) and was evaluated as-is, on principle, without further training.

	v1.1 (ACTIVE)	v1.2 CANDIDATE
no_helmet recall	0.45	0.25
person recall	0.81	0.00

Gate failed on 2 of 3 checks — v1.1 remains active. Nothing overwritten silently.

Active model (v1.1) — held-out Construction-PPE test split

CLASS	PRECISION	RECALL	AP50	AP50-95
person	0.78	0.81	0.83	0.51
helmet	0.92	0.90	0.93	0.52
vest	0.84	0.87	0.88	0.56
no_helmet	0.38	0.17	0.23	0.08

0.559

MAP50

0.276

MAP50-95

LATENCY & VIDEO-EVENT PERFORMANCE

11.1ms

MEDIAN LATENCY (MX450)


88

ACHIEVED FPS

REAL-WORLD INTERVIEW-VIDEO STRESS TEST

Person detected in 354/527 frames, helmet in 143, vest in 401. **no_helmet fired in 0 frames** — a genuine, disclosed weakness (peaks at 0.049 vs. the 0.05 runtime threshold on this real clip).

One screen: live risk, forecast, camera, incidents



- Live gas levels — current CO2 ppm from real ingested readings
- Risk severity — deterministic overall status card (NORMAL / HIGH / CRITICAL...)
- 60-min forecast — physics + GRU hybrid, with uncertainty band
- Time-to-Action — plain-language threshold-crossing estimate
- Camera supervision — live real detections, PPE/zone state per track
- Active incidents — severity, zone, age, one-click Review
- Evidence thumbnails — in the review drawer, real captured frames

Three genuine, real-video-triggered events



PPE COMPLIANT

Helmet + vest detected, confidence 0.78/0.79 — no incident.



PPE_HELMET_OVERHEAD_VIOLATION · HIGH

No positive helmet evidence while inside the overhead-work zone.




PERSON_IN_RESTRICTED_ZONE · HIGH


Foot point dwelled 2.5s inside the configured restricted polygon.

All three are real captured frames (`is_real_camera_frame=true`) from actual incidents opened by the live backend during this session — reviewable end-to-end with JSON/CSV reports in `docs/screenshots/final/07` and `/08`.


Testable without a physical factory



Factory layout: workers, zones, and machinery markers.



Live gas/ventilation sliders and scenario presets.



Click-to-move worker, helmet/vest toggles.

Simulation commands flow through **the exact same ingestion → risk → incident pipeline** a real sensor or camera would use — this is not a decorative mockup, it is a legitimate deterministic test harness (seeded, replayable, used by the automated test suite itself).

The real sequence, start to finish

- 1 Video / sensor input — real camera feed or seeded gas readings
- 2 Prediction / detection — YOLO+ByteTrack or physics+GRU+XGBoost
- 3 Deterministic risk decision — versioned severity policy
- 4 Incident creation — deduplicated, reason-coded, in the database
- 5 Evidence image — real captured frame or labelled schematic
- 6 Manager acknowledgement — Acknowledge → Investigate → Resolve
- 7 Downloadable report — JSON and CSV, per incident

WATCH IT RUN

Full interview-demo recording and annotated video are in the repository:

github.com/M7mdknh/smart-detector

deliverables/Factory_Safety_Sentinel_Interview_Demo.mp4

Verified this session, not stale figures

173

BACKEND TESTS PASSING

18

FRONTEND TESTS PASSING

2/2

PLAYWRIGHT E2E SUITES

0

KNOWN REGRESSIONS

ACCEPTANCE MATRIX (DOCS/ACCEPTANCE_RESULTS.MD)

A01–A18: 17 PASS, 1 PASS WITH LIMITATION, 0 FAIL · E01–E12: 10 PASS, 2 PASS WITH LIMITATION, 0 FAIL

Limitations are disk-quota constraints on a from-scratch vision install and evidence-image honesty labelling — never a functional failure. Docker build/up/down and clean-checkout `make demo` verified end to end in the same audit.

Stated directly, not defensively

- No real factory deployment data — trained/evaluated on construction-site imagery
- Construction-to-factory domain gap is real and measured, not assumed away
- No facial recognition — anonymous, session-local track IDs only, by design
- Evidence images contain worker imagery — retention/privacy policy needed for real use
- The model can and does miss PPE violations (no_helmet is a known weak point)
- **Not certified for real industrial safety decisions**
- A human manager remains responsible for every action taken
- Future work: calibration and domain adaptation on real factory footage

CONCLUSION

A proactive, evidence-based safety system — testable today, extendable tomorrow

- Proactive gas-risk prediction, minutes before a threshold is crossed
- Automated visual supervision of PPE and restricted zones
- Earlier response through a real, reviewable incident workflow
- Evidence-backed, auditable, never a fabricated alert
- Fully testable without a physical factory
- Extendable to real sensors and cameras with no backend rewrite

Questions?